

Contents

Installation of Damaris with Catalyst Support using Spack	1
Install and configure Spack	1
Building Damaris and Paraview Catalyst dependencies.	2
Set the Spack packages.yaml file to specify python3	2
Update the Damaris package file	2
Create a Spack environment	3
Install Damaris	3
Run a Damaris example and have it connect to Paraview	4
Find and set up the example simulation code	4
Launch Paraview GUI	5
Execute the example simulation code	5
Render the simulation data using Paraview	5
Appendix	7
A1. Request resources from Grid5000	7
A2. Set up sshfs	8
A3. Some notes on Graphics drivers on Linux:	8
A4. Build and Install Damaris from git repository	9
A5. Installation of Damaris dependencies on VM with restricted memory	10

Installation of Damaris with Catalyst Support using Spack

This documentation is currently augmenting what is available on <https://project.inria.fr/damaris/>. It is targeted to software versions Damaris 1.3.1, Catalyst 5.6.0 and Spack 0.14.2.

Install and configure Spack

We will be using Spack (<https://spack.io/>) to install the full graph of dependencies of Damaris with Catalyst support. This will require a couple of hours or so to complete and is best done on a compute node with multiple cores. See [Appendix](#) for instructions on how to request extra resources on Grid5000.

1. Obtain Spack

- If Spack is not available (try `which spack`) on your system then download it via `git`

```
# Create a directory for Spack
$ mkdir ~/myspack
$ cd ~/myspack
# download the Spack repository
$ git clone https://github.com/spack/spack.git
$ export SPACK_ROOT=$PWD/spack
$ . spack/share/spack-env.sh
$ spack bootstrap
```

- If it is available you will need to make a copy so that we can change some Spack files to suit our requirements

```
# Create a directory for Spack
$ mkdir ~/myspack
$ spack clone ~/myspack
```

2. Ensure the Spack environment is set up in future shells by adding the following to your `~/.bashrc` file

```
$ cat ~/.bashrc
# Set up your editor of choice here
export EDITOR=nano
export SPACK_ROOT=~/.myspack/spack
. $SPACK_ROOT/share/spack/setup-env.sh
```

3. If running MPI jobs across multiple nodes then create a ~/.profile file that will source the ~/.bashrc file. This will make sure that the same environment is used on all the nodes in login and non-login shells that are created by OpenMPI/MPI.

```
$ cat ~/.profile
...
source ~/.bashrc
...
```

4. Check that your Spack build will use the desired number of cores

```
# Check the config.yaml file
$ cat ~/.spack/config.yaml
config:
  build_stage:
    - ~/.spack/stage

  build_jobs: 4
```

Building Damaris and Paraview Catalyst dependencies.

For Catalyst installation there is a Spack configuration requirement that needs to be made to ensure Python3 is preferred over Python2 for the whole set of dependencies. The Damaris Spack package file may need to be modified to accommodate the Python3 dependency and also to add a CMake flag so that example code is compiled.

Set the Spack packages.yaml file to specify python3

This should prevent installed libraries from using Python2 and Spack *py-packages* seem to get mixed up due to their Python2/3 dependencies.

```
$ cat ~/.spack/packages.yaml
---
packages:
  python:
    version: [3,2]
```

Update the Damaris package file

Make sure the Damaris Spack ‘package.py’ file needs to be updated to specify `catalyst+python3` *not* `catalyst+python`. Also needed in the file are commands to add a CMake `examples` variant and the CMake flag to build the Damaris examples.

An updated Damaris Spack 'package.py' file is available from the [gitlab repository](#) and can be used by setting up a custom Spack repository by adding the path to the `build/spack/repo` directory to the file `~/.spack/repos.yaml`.

```
$ cat ~/.spack/repos.yaml
repos:
- <git repo dir>/damaris/build/spack/repo

$ cat <git repo dir>/damaris/build/spack/repo/repos.yaml
repo:
  namespace: 'damaris.gitlab.dev'
```

All future `spack install damaris` commands will preferentially use the custom repo, unless explicitly overridden using a specific namespace e.g. `spack install builtin.damaris`. For further documentation on Spack repositories please see [Spack.io Repositories](#)

The differences between the package.py code that is obtained from Spack and that which is present from the Damaris gitlab code repository are something like as follows:

```
>spack edit damaris.gitlab.dev.damaris
...
variant('examples', default=False, description='Enables compilation and installation of
the examples code')
...
depends_on('catalyst+python3', when='+catalyst@5.6')
...
cmake_args()
...
if (self.spec.variants['examples'].value):
    args.extend(['-DENABLE_EXAMPLES:BOOL=ON'])
```

Create a Spack environment

We will create a Spack *environment* to work from. This is a filesystem area that will contain all the libraries needed to run and build Damaris and its dependent executable simulation code. Environments are easily created from Spack installed libraries and can be removed easily without removing the underlying Spack installed library. The activation of the environment will set the `PATH`, `LIBRARY_PATH`, `CPATH`, `LD_LIBRARY_PATH`, `PKG_CONFIG_PATH`, `MANPATH` environment variables to the view of the system specific to the installed libraries and dependencies. One should note that common tools may not be available without installing them specifically in the environment (e.g. the `nano` editor) or possibly loading them using modules.

```
# create and the activate a spack environment
$ spack env create damaris-catalyst
$ spack env activate damaris-catalyst
```

To support OpenMPI/MPI ensure the Spack environment is loaded by your shell by adding the following to your `~/.bashrc` file

```
$ cat ~/.bashrc
...
spack env activate damaris-catalyst
```

Install Damaris

The following `spack install` command should build and install Damaris and all of its dependencies including the Paraview Catalyst libraries. This process can take a long time (multiple hours). If a particular part fails ([LLVM for example](#)) then the Spack environment should be removed using `spack env rm <env-name>` and then re-created. The `--keep-stage` option is needed so that the Damaris examples are not removed when the build directory is cleaned, as the Damaris Spack install does not install the examples for us.

```
# Install your favourite editor
$ spack install nano
# Install dependencies without keeping the staging directory (saves disk space)
$ spack install catalyst@5.6.0+osmesa+rendering ^mesa+osmesa swr=avx ^llvm@8.0.0 target=
  sandybridge

# Install Damaris and dependencies
$ spack install --keep-stage damaris+catalyst+hdf5+examples+fortran
```

The Mesa ‘swr’ option can be specified for x86_64 CPUs dependent on the architecture of the CPU on the compute nodes. I have specified `swr=avx` and `target=sandybridge` so that libs are compatible with my VM CPU.

Run a Damaris example and have it connect to Paraview

Download and install a Paraview version that matches the version of Catalyst that was installed via Spack available here: <https://www.paraview.org/download/>. There are pre-compiled executables available for Windows, Mac and Linux. When using Linux on compute nodes or on a VM no GPU support then check out notes on running with Mesa in [appendix below](#).

Find and set up the example simulation code

To access the examples from the Damaris Spack build directory use the `spack cd` command and navigate to the build directory.

```
# Check what the Spack install package configuration was
# (from within the environment)
$ spack config edit

# 'spack cd' gets us to the install directory
$ spack cd damaris+catalyst+hdf5+examples+fortran ^catalyst@5.6.0+osmesa+rendering ^mesa+
  osmesa swr=avx ^llvm@8.0.0 target=sandybridge

# Now find the build directory and examples that we kept
# using '--keep-stage'
$ cd ../spack-build/examples/paraview
# Copy the input damaris xml files and catalyst python scripts
# to the executables
$ cp ../../../../spack-src/examples/paraview/* .
```

Now we have the examples compiled, check the configuration of the xml file

1. Select the number of dedicated cores (per node).
2. Ensure image.xml file has the current path to the `image.py` Catalyst script corrected
3. Ensure the `<parameter name="size" value = (mpi_proc_per_node - dedicated cores)‘`

```

<architecture>
...
<dedicated cores="1" nodes="0" />
...
</architecture>
<data>
...
<parameter name="size" type="int" value="3" />
...
</data>
...
<paraview>
  <script>image.py</script>
</paraview>

```

Launch Paraview GUI

Now launch the Paraview GUI and then select the Catalyst menu and select “Connect”




Figure 1: Paraview Catalyst toolbar

In the text box that is presented you can select the port on which to listen, which will need to be the same one as specified in the example code Catalyst .py file python script (the .py script is specified in the Damaris .xml configuration file in the <paraview tag). Port 22222 is a default value so should work as given, so now click OK.




Figure 2: Paraview Catalyst Connect box

Execute the example simulation code

The next step is to execute the simulation code using `mpirun` and the appropriate number of processes. For testing on a single node, the OpenMPI `--oversubscribe` flag may be needed if there are not as many cores available as processes being run.

```

# Make sure you are in the spack-build/examples/paraview
# examples directory (use `spack cd ...`)
# This is a typical Damaris example where the Damaris

```

```
# configuration xml file is specified on the command line
$ mpirun --oversubscribe -np 4 ./image "image.xml"
```

Render the simulation data using Paraview

ParaView should soon find the connection and present a *catalyst* option in the *Pipeline Browser* tab




Figure 3: Paraview GUI: pipeline browser

The data in this case is named *input* and can be selected for display within the *builtin* pipeline viewer and then shown with default bounds and color by selecting the *eye* icon to view:




Figure 4: Paraview GUI: Select the input builtin->dataset




Figure 5: Paraview GUI: Selecting the eye icon to render the data

Now that the data is available in the Pipeline Browser the full set of visualization possibilities are available through ParaView. To start with the data is rendered in a very conservative manner - as a bounded box with single set color. The dataset can be rendered and colored by modifying the *Representation* and *Color* of the dataset using either the *Active variable control* toolbar or the *Properties* tab on the *Pipeline Browser* View panel as seen in fig.6 and fig.7 below.



Figure 6: Paraview GUI: Active variable control bar

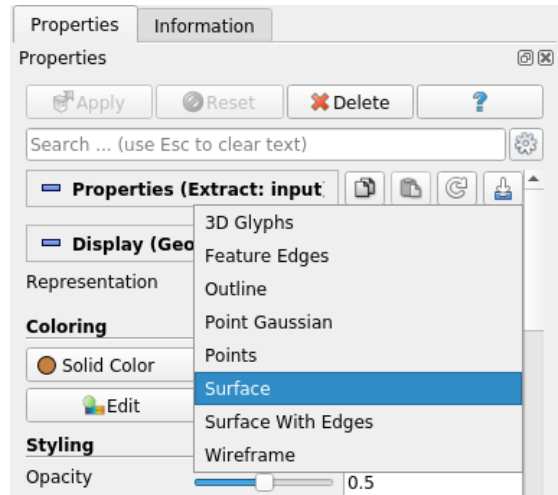


Figure 7: Paraview GUI: Properties panel. Choose a representation and select a Coloring method that highlights a variable of interest

Appendix

A1. Request resources from Grid5000

[Back](#)

```
$ ssh <username>@rennes

# Check if there are nodes available
$ funk -m date -r rennes -w 04:00:00

# request 4 cores (or more)
$ oarsub -I -p "cluster='paravance'" -l /nodes=1/core=16,walltime=04:00:00

# Check that your spack build will use all cores:
$ cat ~/.spack/config.yaml
config:
build_stage:
- ~/.spack/stage

build_jobs: 16
```

We will also need OpenMPI set up correctly to use the `oarsh` helper to talk to nodes

```
$ cat ~/.openmpi/mca-params.conf
# orte_rsh_agent=oarsh # for pre-ompi@3.1
plm_rsh_agent=oarsh
filem_rsh_agent=oarcp
```

The `oarsh` command is the same one used to access nodes requested via `oarsub`. This is how you use `oarsh` to log into the job running your node

```

# check what job I am running (if forgotten)
$ oarstat -cu <username>
Job id      Name      User      Submission Date      S Queue
-----
1270759      <username>      2020-05-27 15:43:50 R default

# export the job number of interest
$ OAR_JOB_ID=1270759
# print the given nodes provided by the job
$ oarstat -fj $OAR_JOB_ID | grep assigned_hostnames | sed -n 's/^.*= //p'
paravance-28.rennes.grid5000.fr

# the OAR_JOB_ID variable is used by oarsh
$ oarsh paravance-28
# you should now be on the node

```

A2. Set up sshfs

Set up sshfs between a VM and Grid5000 user Rennes NFS filesystem

```

$ ssh <vm-name>.grid5000.fr
$ sudo apt-get update
$ sudo apt-get install sshfs
$ sudo gpasswd -a <username> fuse # (add myself to fuse group)

# make the local shared directory
$ mkdir ~/sshfslocal
# Make the remote shared directory
$ ssh <username>@frennes.rennes.grid5000.fr mkdir /home/<username>/sshfsshare
# Set up the filesystem between the two directories
$ sshfs <username>@frennes.rennes.grid5000.fr:/home/<username>/sshfsshare /home/<username>
  >/sshfslocal

```

A3. Some notes on Graphics drivers on Linux:

Return to Run Damaris Example

Running Catalyst on systems without dedicated graphics hardware (i.e. Virtual Machines and compute nodes of clusters) requires Mesa 3D graphics libraries which are used for OpenGL rendering support. Mesa uses the llvm infrastructure so llvm will be installed as a dependency (see [LLVM Installation Issues](#)). For a better understanding of the Mesa driver infrastructure and its multiple options, please see: www.paraview.org/ParaView_And_Mesa_3D)

OSMesa (“Off Screen Mesa”) is the front end to various drivers that provide different levels of acceleration support. For X86_64 with AVX (or greater) instructions there is the Gallium driver *swr* otherwise *llvmpipe* is available for multiple CPU targets. There is also a fallback single threaded driver *softpipe*.

Once LLVM and Mesa are installed there are some environment variables to tweak.
(see: www.mesa3d.org/envvars.html)

```

$ export GALLIUM_DRIVER=softpipe|llvmpipe|swr

# If llvmpipe is chosen, the number of threads to use is

```



```

# selected as follows
# llmpipe is only threaded in the pixel operations
# so the thread level has no improvement in vertex rendering.
# N.B. Catalyst/MPI adds capability for vertex rendering.
$ export LP_NUM_THREADS=2

# if sur is available (X86_64 only), it is threaded in vertex and pixel operations
$ export KNOB_MAX_WORKER_THREADS=1..256

```

Other performance tips for using the Mesa are available www.mesa3d.org/perf.html

A4. Build and Install Damaris from git repository

Use the following guide to manually compile a development version of Damaris within a Spack environment that has Catalyst installed.

Download development version from gitlab

```

# N.B. git may not be available from the spack environemnt (i.e. first use despactivate)
$ git clone https://gitlab.inria.fr/Damaris/damaris-development.git

```

Build Damaris examples with Catalyst support (no VisIt)

N.B. `install_path` can be determined using `spack env st`

```

$ spack env activate damaris-catalyst

$ cd damaris-development/build
$ mkdir mybuild
$ cd mybuild

$ export install_path=~/.spack/var/spack/environments/damaris-catalyst/.spack-env/view
$ hdf5_arg="-DENABLE_HDF5=ON -DHDF5_ROOT=$install_path"
$ visit_arg="-DENABLE_VISIT=OFF -DVisIt_ROOT=$install_path"
$ catalyst_arg="-DENABLE_CATALYST=ON -DParaView_DIR=$install_path"
$ cmake ../../.. -DCMAKE_INSTALL_PREFIX:PATH=$install_path \
-DBOOST_ROOT=$install_path \
-DXSD_ROOT=$install_path \
-DXercesC_ROOT=$install_path \
-DCppUnit_ROOT=$install_path \
-DCMAKE_CXX_COMPILER=mpicxx \
-DCMAKE_C_COMPILER=mpicc \
-DENABLE_TESTS=OFF \
-DENABLE_EXAMPLES=ON \
-DBUILD_SHARED_LIBS=ON \
$visit_arg \
$hdf5_arg \
$catalyst_arg

$ cd examples/paraview
$ make -j 4
$ make install

```

A5. Installation of Damaris dependencies on VM with restricted memory

There can be multiple issues found when installing Catalyst on systems with limited amount of available RAM (e.g. 2GB running XFCE over Xrdp/VNC). This will affect LLVM and Catalyst builds (vtkDataArray.cxx compile was a common fail point). On machines where LLVM did compile then the Mesa build would fail for LLVM versions ≥ 9.0 . The Spack Mesa installation is restricted to versions $\leq \text{mesa@18.3.6}$ when installing with Spack as the Mesa build system was changed to use meson which is not yet supported by Spack.

To overcome LLVM installation issues

LLVM was a source of multiple issues during installation of Damaris with visualization support when using Spack. 1. Require a machine with $>2\text{GB}$ RAM otherwise the LLVM build may fail (see below). 2. LLVM versions $\leq 8.0.0$ are compatible with the Mesa installed by Spack (mesa@18.3.6)

A way around the RAM requirement is to get Spack to use a system installed llvm. You may need to install clang/llvm using your system package manager first.

```
# use sudo-g5k for Grid5000 builds
$ sudo apt-get install clang
$ clang --version

# Add dependency to packages.yaml
$ cat ~/.spack/packages.yaml
---
packages:
  python:
    version: [3,2]
  llvm:
    buildable: False
    paths:
      llvm@7.0.0: /usr/lib/llvm-7
    version: [7.0]
```

Fix for Catalyst build fail

The Catalyst libraries were also found to require as much as available of 2GB RAM for building. Logging out of the VNC session and killing the server of my build machine and compilation from a terminal without TigerVNC and XFCE overheads succeeded, albeit with the system installed llvm.